

DataPreprocessing-EDA

April 13, 2025

Original Data source <https://nihcc.app.box.com/v/ChestXray-NIHCC>

Google Healthcare APIs <https://cloud.google.com/healthcare-api/docs/resources/public-datasets/nih-chest>

```
[1]: !pip install kagglehub  
      !pip install kagglehub[pandas-datasets]  
      !pip install wget  
      !pip install keras-tuner
```

```
Requirement already satisfied: kagglehub in /usr/local/lib/python3.11/dist-packages (0.3.11)  
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from kagglehub) (24.2)  
Requirement already satisfied: pyyaml in /usr/local/lib/python3.11/dist-packages (from kagglehub) (6.0.2)  
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from kagglehub) (2.32.3)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from kagglehub) (4.67.1)  
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->kagglehub) (3.4.1)  
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->kagglehub) (3.10)  
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->kagglehub) (2.3.0)  
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->kagglehub) (2025.1.31)  
Requirement already satisfied: kagglehub[pandas-datasets] in /usr/local/lib/python3.11/dist-packages (0.3.11)  
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-packages (from kagglehub[pandas-datasets]) (24.2)  
Requirement already satisfied: pyyaml in /usr/local/lib/python3.11/dist-packages (from kagglehub[pandas-datasets]) (6.0.2)  
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-packages (from kagglehub[pandas-datasets]) (2.32.3)  
Requirement already satisfied: tqdm in /usr/local/lib/python3.11/dist-packages (from kagglehub[pandas-datasets]) (4.67.1)  
Requirement already satisfied: pandas in /usr/local/lib/python3.11/dist-packages
```

```

(from kagglehub[pandas-datasets]) (2.2.2)
Requirement already satisfied: numpy>=1.23.2 in /usr/local/lib/python3.11/dist-
packages (from pandas->kagglehub[pandas-datasets]) (2.0.2)
Requirement already satisfied: python-dateutil>=2.8.2 in
/usr/local/lib/python3.11/dist-packages (from pandas->kagglehub[pandas-
datasets]) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-
packages (from pandas->kagglehub[pandas-datasets]) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-
packages (from pandas->kagglehub[pandas-datasets]) (2025.2)
Requirement already satisfied: charset-normalizer<4,>=2 in
/usr/local/lib/python3.11/dist-packages (from requests->kagglehub[pandas-
datasets]) (3.4.1)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-
packages (from requests->kagglehub[pandas-datasets]) (3.10)
Requirement already satisfied: urllib3<3,>=1.21.1 in
/usr/local/lib/python3.11/dist-packages (from requests->kagglehub[pandas-
datasets]) (2.3.0)
Requirement already satisfied: certifi>=2017.4.17 in
/usr/local/lib/python3.11/dist-packages (from requests->kagglehub[pandas-
datasets]) (2025.1.31)
Requirement already satisfied: six>=1.5 in /usr/local/lib/python3.11/dist-
packages (from python-dateutil>=2.8.2->pandas->kagglehub[pandas-datasets])
(1.17.0)
Collecting wget
  Downloading wget-3.2.zip (10 kB)
  Preparing metadata (setup.py) ... done
Building wheels for collected packages: wget
  Building wheel for wget (setup.py) ... done
  Created wheel for wget: filename=wget-3.2-py3-none-any.whl size=9655
sha256=018f394ef0d70790c422fee28405aa5b022ca62df5c6f71b9b76aa04822c2938
  Stored in directory: /root/.cache/pip/wheels/40/b3/0f/a40dbd1c6861731779f62cc4
babcb234387e11d697df70ee97
Successfully built wget
Installing collected packages: wget
Successfully installed wget-3.2
Collecting keras-tuner
  Downloading keras_tuner-1.4.7-py3-none-any.whl.metadata (5.4 kB)
Requirement already satisfied: keras in /usr/local/lib/python3.11/dist-packages
(from keras-tuner) (3.8.0)
Requirement already satisfied: packaging in /usr/local/lib/python3.11/dist-
packages (from keras-tuner) (24.2)
Requirement already satisfied: requests in /usr/local/lib/python3.11/dist-
packages (from keras-tuner) (2.32.3)
Collecting kt-legacy (from keras-tuner)
  Downloading kt_legacy-1.0.5-py3-none-any.whl.metadata (221 bytes)
Requirement already satisfied: absl-py in /usr/local/lib/python3.11/dist-
packages (from keras->keras-tuner) (1.4.0)

```

Requirement already satisfied: numpy in /usr/local/lib/python3.11/dist-packages (from keras->keras-tuner) (2.0.2)

Requirement already satisfied: rich in /usr/local/lib/python3.11/dist-packages (from keras->keras-tuner) (13.9.4)

Requirement already satisfied: namex in /usr/local/lib/python3.11/dist-packages (from keras->keras-tuner) (0.0.8)

Requirement already satisfied: h5py in /usr/local/lib/python3.11/dist-packages (from keras->keras-tuner) (3.13.0)

Requirement already satisfied: optree in /usr/local/lib/python3.11/dist-packages (from keras->keras-tuner) (0.14.1)

Requirement already satisfied: ml-dtypes in /usr/local/lib/python3.11/dist-packages (from keras->keras-tuner) (0.4.1)

Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.11/dist-packages (from requests->keras-tuner) (3.4.1)

Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.11/dist-packages (from requests->keras-tuner) (3.10)

Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.11/dist-packages (from requests->keras-tuner) (2.3.0)

Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.11/dist-packages (from requests->keras-tuner) (2025.1.31)

Requirement already satisfied: typing-extensions>=4.5.0 in /usr/local/lib/python3.11/dist-packages (from optree->keras->keras-tuner) (4.13.1)

Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras->keras-tuner) (3.0.0)

Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.11/dist-packages (from rich->keras->keras-tuner) (2.18.0)

Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.11/dist-packages (from markdown-it-py>=2.2.0->rich->keras->keras-tuner) (0.1.2)

Downloading keras_tuner-1.4.7-py3-none-any.whl (129 kB)

129.1/129.1 kB

3.4 MB/s eta 0:00:00

Downloading kt_legacy-1.0.5-py3-none-any.whl (9.6 kB)

Installing collected packages: kt-legacy, keras-tuner

Successfully installed keras-tuner-1.4.7 kt-legacy-1.0.5

0.1 Load Libraries

```
[2]: import os
import zipfile
import seaborn as sns
import numpy as np
import kagglehub
from kagglehub import KaggleDatasetAdapter
import pandas as pd
import matplotlib.pyplot as plt
import cv2
```

```
import urllib.request

import tensorflow as tf
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.model_selection import train_test_split
from sklearn.utils.class_weight import compute_class_weight
from tensorflow.keras import layers, models, Input, Model
from kerastuner import HyperModel
from kerastuner.tuners import RandomSearch
```

<ipython-input-2-b7ab669529a1>:17: DeprecationWarning: `import kerastuner` is deprecated, please use `import keras_tuner`.

```
from kerastuner import HyperModel
```

```
[3]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[4]: # Global flags
SKIP_BOUNDING_BOX = True
SKIP_DOWNLOAD = False
SKIP_UNZIP = False

DRIVE_PATH = "/content/drive/MyDrive/AAI-590_Collabs"
RESIZED_IMAGES_ZIP_PATH = "/content/drive/MyDrive/AAI-590_Collabs"
RESIZED_IMAGES_PATH = "/content/images_resized/images_resized";
```

```
[5]: SKIP_DOWNLOAD = os.path.exists(RESIZED_IMAGES_ZIP_PATH)
SKIP_UNZIP = os.path.exists(RESIZED_IMAGES_PATH)
```

```
[6]: # print current variables
print("SKIP_DOWNLOAD: ", SKIP_DOWNLOAD)
print("SKIP_UNZIP: ", SKIP_UNZIP)
```

SKIP_DOWNLOAD: True

SKIP_UNZIP: False

0.2 Load Dataset

```
[7]: # Set the dataset path
dataset_name = "nih-chest-xrays/data"
version = 3
# Set the path to the file you'd like to load
file_path = "Data_Entry_2017.csv"
file_path_bbox = "BBBox_List_2017.csv"

gcloud_url_base = 'https://storage.googleapis.com/
↳gcs-public-data--healthcare-nih-chest-xray/png/'
```

```
[8]: # Load the latest version
df = kagglehub.load_dataset(
    KaggleDatasetAdapter.PANDAS,
    dataset_name,
    file_path,
    # Provide any additional arguments like
    # sql_query or pandas_kwargs. See the
    # documenation for more information:
    # https://github.com/Kaggle/kagglehub/blob/main/README.
    ↪md#kaggledatasetadapterpandas
)

df_box_list = kagglehub.load_dataset(
    KaggleDatasetAdapter.PANDAS,
    dataset_name,
    file_path_bbox
)
```

<ipython-input-8-1e51267dc0e1>:2: DeprecationWarning: load_dataset is deprecated and will be removed in future version.

```
df = kagglehub.load_dataset(

Downloading from https://www.kaggle.com/api/v1/datasets/download/nih-chest-
xrays/data?dataset_version_number=3&file_name=Data_Entry_2017.csv...

100%|          | 924k/924k [00:00<00:00, 101MB/s]

Extracting zip of Data_Entry_2017.csv...
```

<ipython-input-8-1e51267dc0e1>:12: DeprecationWarning: load_dataset is deprecated and will be removed in future version.

```
df_box_list = kagglehub.load_dataset(

Downloading from https://www.kaggle.com/api/v1/datasets/download/nih-chest-
xrays/data?dataset_version_number=3&file_name=BBox_List_2017.csv...

100%|          | 90.2k/90.2k [00:00<00:00, 2.13MB/s]
```

```
[9]: print(df['View Position'].value_counts())
```

```
View Position
PA      67310
AP      44810
Name: count, dtype: int64
```

```
[10]: # keep original dataframe for reference
df_locked = df.copy()
```

0.3 Remove all where “View Position” column value is “AP”

AP means “anteroposterior dimension” which is an X-ray from front-to-back This wil affect the training with both back-to-front and front-to-back images of MRIs

```
[11]: # Entries before removal
print(f"Before 'AP' removal: {df['View Position'].value_counts()}")

# Entries after removal
df = df[df['View Position'] != 'AP']

# Remaining data is 66.57% of total initial data
print(f"After 'AP' removal: {df['View Position'].value_counts()}")
```

```
Before 'AP' removal: View Position
PA      67310
AP      44810
Name: count, dtype: int64
After 'AP' removal: View Position
PA      67310
Name: count, dtype: int64
```

```
[12]: links = [
    "https://nihcc.box.com/shared/static/vfk49d74nhbxq3nqjg0900w5nvkorp5c.gz",
    "https://nihcc.box.com/shared/static/i28rlmbvmfjbl8p2n3ril0pptcmcu9d1.gz",
    "https://nihcc.box.com/shared/static/f1t00wrtdk94satdfb9olcolqx20z2jp.gz",
    "https://nihcc.box.com/shared/static/0aowwzs5lhjrceb3qp67ahp0rd1l1etg.gz",
    "https://nihcc.box.com/shared/static/v5e3goj22zr6h8tzualxfsqlqaygfbsn.gz",
    "https://nihcc.box.com/shared/static/asi7ikud9jwnkrnkj99jnpfkjdes7l6l.gz",
    "https://nihcc.box.com/shared/static/jn1b4mw4n6lnh74ovmcjb8y48h8xj07n.gz",
    "https://nihcc.box.com/shared/static/tvpxmn7qyrgl0w8wfh9kqfjskv6nmm1j.gz",
    "https://nihcc.box.com/shared/static/upyy3ml7qdumlgk2rfcvlb9k6gvqq2pj.gz",
    "https://nihcc.box.com/shared/static/l6nilvfa9cg3s28tqv1qc1olm3gnz54p.gz",
    "https://nihcc.box.com/shared/static/hhq8fkdgvvari67vfhs7ppg2w6ni4jze.gz",
    "https://nihcc.box.com/shared/static/ioqwi20ihqwyr8pf4c24eazhh281pbu.gz",
]
```

```
[13]: # Create a dictionary for folder locations
folder_ranges = {
    "images_001": (0, 4998), # Adjusted to 0-based index
    "images_002": (4999, 14998),
    "images_003": (14999, 24998),
    "images_004": (24999, 34998),
    "images_005": (34999, 44998),
    "images_006": (44999, 54998),
    "images_007": (54999, 64998),
    "images_008": (64999, 74998),
    "images_009": (74999, 84998),
}
```

```

    "images_010": (84999, 94998),
    "images_011": (94999, 104998),
    "images_012": (104999, 112120)
}

def get_image_folder(df, image_name):
    if image_name in df["Image Index"].values:
        image_index = df[df["Image Index"] == image_name].index[0] # Get row index
        # print(f"Image {image_name} is at index {image_index}") # Debugging output

        for folder, (start, end) in folder_ranges.items():
            if start <= image_index <= end:
                return folder

    return None # If not found

```

```

[14]: display(df.head())
      display(df.tail())
      display(df.columns)

```

	Image Index	Finding Labels	Follow-up #	Patient ID	\
0	00000001_000.png	Cardiomegaly	0	1	
1	00000001_001.png	Cardiomegaly Emphysema	1	1	
2	00000001_002.png	Cardiomegaly Effusion	2	1	
3	00000002_000.png	No Finding	0	2	
4	00000003_000.png	Hernia	0	3	

	Patient Age	Patient Gender	View Position	OriginalImage[Width	Height]	\
0	58	M	PA	2682	2749	
1	58	M	PA	2894	2729	
2	58	M	PA	2500	2048	
3	81	M	PA	2500	2048	
4	81	F	PA	2582	2991	

	OriginalImagePixelSpacing[x	y]	Unnamed: 11
0	0.143	0.143	NaN
1	0.143	0.143	NaN
2	0.168	0.168	NaN
3	0.171	0.171	NaN
4	0.143	0.143	NaN

	Image Index	Finding Labels	Follow-up #	Patient ID	\
112115	00030801_001.png	Mass Pneumonia	1	30801	
112116	00030802_000.png	No Finding	0	30802	
112117	00030803_000.png	No Finding	0	30803	
112118	00030804_000.png	No Finding	0	30804	

112119	00030805_000.png	No Finding	0	30805
--------	------------------	------------	---	-------

	Patient	Age	Patient	Gender	View	Position	OriginalImage[Width	\
112115		39		M		PA	2048	
112116		29		M		PA	2048	
112117		42		F		PA	2048	
112118		30		F		PA	2048	
112119		27		M		PA	2048	

	Height]	OriginalImagePixelSpacing[x	y]	Unnamed: 11
112115	2500	0.168	0.168	NaN
112116	2500	0.168	0.168	NaN
112117	2500	0.168	0.168	NaN
112118	2500	0.168	0.168	NaN
112119	2500	0.171	0.171	NaN

```
Index(['Image Index', 'Finding Labels', 'Follow-up #', 'Patient ID',
      'Patient Age', 'Patient Gender', 'View Position', 'OriginalImage[Width',
      'Height]', 'OriginalImagePixelSpacing[x', 'y]', 'Unnamed: 11'],
      dtype='object')
```

0.4 We want to have 7 generalized classes from the original 15

Take values from “Finding Labels” and convert them into more generalized labels

```
[15]: # Create a list to store all unique labels
all_labels = []

# Iterate over the 'Finding Labels' column
for index, row in df.iterrows():
    labels = row['Finding Labels'].split('|')
    for label in labels:
        all_labels.append(label)
```

```
# Get unique labels and print them
all_labels = list(set(all_labels))
print(f"All possible options in 'Finding Labels': {all_labels}")
```

```
All possible options in 'Finding Labels': ['Effusion', 'Pneumonia', 'Hernia',
'Infiltration', 'Atelectasis', 'Mass', 'Pneumothorax', 'Emphysema',
'Pleural Thickening', 'No Finding', 'Consolidation', 'Fibrosis', 'Edema',
'Cardiomegaly', 'Nodule']
```

```
[16]: def generalize_labels(label):
    if label in ['Pneumonia', 'Consolidation', 'Infiltration']:
        return 'Infection/Infiltration'
    elif label in ['Edema', 'Effusion', 'Pleural Thickening']:
        return 'Fluid Related Issues'
    elif label in ['Atelectasis', 'Pneumothorax', 'Fibrosis', 'Emphysema']:
```



```

        return 'Lung Structure Issues'
    elif label in ['Nodule', 'Mass']:
        return 'Nodule/Mass'
    elif label == 'Cardiomegaly':
        return 'Cardiac Issues'
    elif label == 'Hernia':
        return 'Hernia'
    else:
        return label # If we don't detect an issue 'No Finding'

df['Finding Labels'] = df['Finding Labels'].apply(lambda x: '|'.
    ↪join([generalize_labels(label) for label in x.split('|')]))

# Example:
display(df.head()) # View the updated DataFrame

```

	Image Index	Finding Labels	Follow-up #	\
0	00000001_000.png	Cardiac Issues	0	
1	00000001_001.png	Cardiac Issues Lung Structure Issues	1	
2	00000001_002.png	Cardiac Issues Fluid Related Issues	2	
3	00000002_000.png	No Finding	0	
4	00000003_000.png	Hernia	0	

	Patient ID	Patient Age	Patient Gender	View Position	OriginalImage[Width	\
0	1	58	M	PA	2682	
1	1	58	M	PA	2894	
2	1	58	M	PA	2500	
3	2	81	M	PA	2500	
4	3	81	F	PA	2582	

	Height]	OriginalImagePixelSpacing[x	y]	Unnamed: 11
0	2749	0.143	0.143	NaN
1	2729	0.143	0.143	NaN
2	2048	0.168	0.168	NaN
3	2048	0.171	0.171	NaN
4	2991	0.143	0.143	NaN

```

[17]: display(df.head())
      display(df.tail())
      display(df.columns)

```

	Image Index	Finding Labels	Follow-up #	\
0	00000001_000.png	Cardiac Issues	0	
1	00000001_001.png	Cardiac Issues Lung Structure Issues	1	
2	00000001_002.png	Cardiac Issues Fluid Related Issues	2	
3	00000002_000.png	No Finding	0	
4	00000003_000.png	Hernia	0	

	Patient ID	Patient Age	Patient Gender	View Position	OriginalImage[Width \
0	1	58	M	PA	2682
1	1	58	M	PA	2894
2	1	58	M	PA	2500
3	2	81	M	PA	2500
4	3	81	F	PA	2582

	Height]	OriginalImagePixelSpacing[x y]	Unnamed: 11
0	2749	0.143 0.143	NaN
1	2729	0.143 0.143	NaN
2	2048	0.168 0.168	NaN
3	2048	0.171 0.171	NaN
4	2991	0.143 0.143	NaN

	Image Index	Finding Labels	Follow-up # \
112115	00030801_001.png	Nodule/Mass Infection/Infiltration	1
112116	00030802_000.png	No Finding	0
112117	00030803_000.png	No Finding	0
112118	00030804_000.png	No Finding	0
112119	00030805_000.png	No Finding	0

	Patient ID	Patient Age	Patient Gender	View Position \
112115	30801	39	M	PA
112116	30802	29	M	PA
112117	30803	42	F	PA
112118	30804	30	F	PA
112119	30805	27	M	PA

	OriginalImage[Width	Height]	OriginalImagePixelSpacing[x y]	\
112115	2048	2500	0.168 0.168	
112116	2048	2500	0.168 0.168	
112117	2048	2500	0.168 0.168	
112118	2048	2500	0.168 0.168	
112119	2048	2500	0.171 0.171	

	Unnamed: 11
112115	NaN
112116	NaN
112117	NaN
112118	NaN
112119	NaN

```
Index(['Image Index', 'Finding Labels', 'Follow-up #', 'Patient ID',
      'Patient Age', 'Patient Gender', 'View Position', 'OriginalImage[Width',
      'Height]', 'OriginalImagePixelSpacing[x', 'y]', 'Unnamed: 11'],
      dtype='object')
```

```
[18]: display(df.describe())
display(df.info())
```

	Follow-up #	Patient ID	Patient Age	OriginalImage[Width \	
count	67310.000000	67310.000000	67310.000000	67310.000000	
mean	4.786317	14396.542802	47.352979	2632.590016	
std	9.403191	8559.885944	16.289550	374.573816	
min	0.000000	1.000000	1.000000	1143.000000	
25%	0.000000	7157.250000	36.000000	2500.000000	
50%	1.000000	14112.000000	49.000000	2678.000000	
75%	5.000000	21117.750000	59.000000	2992.000000	
max	156.000000	30805.000000	412.000000	3056.000000	

	Height]	OriginalImagePixelSpacing[x	y]	Unnamed: 11
count	67310.000000	67310.000000	67310.000000	0.0
mean	2652.208468	0.153868	0.153868	NaN
std	396.607849	0.017179	0.017179	NaN
min	1001.000000	0.115000	0.115000	NaN
25%	2411.000000	0.143000	0.143000	NaN
50%	2885.000000	0.143000	0.143000	NaN
75%	2991.000000	0.168000	0.168000	NaN
max	3056.000000	0.194336	0.194336	NaN

```
<class 'pandas.core.frame.DataFrame'>
```

```
Index: 67310 entries, 0 to 112119
```

```
Data columns (total 12 columns):
```

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	Image Index	67310 non-null	object
1	Finding Labels	67310 non-null	object
2	Follow-up #	67310 non-null	int64
3	Patient ID	67310 non-null	int64
4	Patient Age	67310 non-null	int64
5	Patient Gender	67310 non-null	object
6	View Position	67310 non-null	object
7	OriginalImage[Width	67310 non-null	int64
8	Height]	67310 non-null	int64
9	OriginalImagePixelSpacing[x	67310 non-null	float64
10	y]	67310 non-null	float64
11	Unnamed: 11	0 non-null	float64

```
dtypes: float64(3), int64(5), object(4)
```

```
memory usage: 6.7+ MB
```

```
None
```

```
[19]: display(df_box_list.head())
display(df_box_list.describe())
display(df_box_list.info())
```

	Image Index	Finding Label	Bbox [x	y	w	\
0	00013118_008.png	Atelectasis	225.084746	547.019217	86.779661	
1	00014716_007.png	Atelectasis	686.101695	131.543498	185.491525	
2	00029817_009.png	Atelectasis	221.830508	317.053115	155.118644	
3	00014687_001.png	Atelectasis	726.237288	494.951420	141.016949	
4	00017877_001.png	Atelectasis	660.067797	569.780787	200.677966	

	h]	Unnamed: 6	Unnamed: 7	Unnamed: 8
0	79.186441	NaN	NaN	NaN
1	313.491525	NaN	NaN	NaN
2	216.949153	NaN	NaN	NaN
3	55.322034	NaN	NaN	NaN
4	78.101695	NaN	NaN	NaN

	Bbox [x	y	w	h]	Unnamed: 6	Unnamed: 7	\
count	984.000000	984.000000	984.000000	984.000000	0.0	0.0	
mean	398.806111	405.425364	256.334708	252.302547	NaN	NaN	
std	222.700868	166.309995	167.629620	159.443635	NaN	NaN	
min	5.417989	12.837934	27.306667	21.617778	NaN	NaN	
25%	203.093333	293.869045	136.533333	115.674074	NaN	NaN	
50%	340.249735	412.850794	214.340942	216.949153	NaN	NaN	
75%	607.959365	521.641995	311.832381	367.902430	NaN	NaN	
max	905.887831	876.980783	901.120000	873.379894	NaN	NaN	

	Unnamed: 8
count	0.0
mean	NaN
std	NaN
min	NaN
25%	NaN
50%	NaN
75%	NaN
max	NaN

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 984 entries, 0 to 983

Data columns (total 9 columns):

#	Column	Non-Null Count	Dtype
0	Image Index	984 non-null	object
1	Finding Label	984 non-null	object
2	Bbox [x	984 non-null	float64
3	y	984 non-null	float64
4	w	984 non-null	float64
5	h]	984 non-null	float64
6	Unnamed: 6	0 non-null	float64
7	Unnamed: 7	0 non-null	float64
8	Unnamed: 8	0 non-null	float64

dtypes: float64(7), object(2)

memory usage: 69.3+ KB

None

```
[20]: # Fix column names
df_box_list = df_box_list.rename(columns={'Bbox [x': 'x', 'h]': 'h'})
df_box_list.head()
```

```
[20]:
```

	Image	Index	Finding	Label	x	y	w	\
0	00013118_008.png		Atelectasis		225.084746	547.019217	86.779661	
1	00014716_007.png		Atelectasis		686.101695	131.543498	185.491525	
2	00029817_009.png		Atelectasis		221.830508	317.053115	155.118644	
3	00014687_001.png		Atelectasis		726.237288	494.951420	141.016949	
4	00017877_001.png		Atelectasis		660.067797	569.780787	200.677966	

	h	Unnamed: 6	Unnamed: 7	Unnamed: 8
0	79.186441	NaN	NaN	NaN
1	313.491525	NaN	NaN	NaN
2	216.949153	NaN	NaN	NaN
3	55.322034	NaN	NaN	NaN
4	78.101695	NaN	NaN	NaN

0.5 EDA

```
[21]: # Analyze class distribution
class_counts = df['Finding Labels'].value_counts()
num_classes = df['Finding Labels'].nunique()
display(f"Number of different classes: {num_classes}")
```

'Number of different classes: 356'

One image can contain one or more labels - thus the number of classes initially shows as 836.

```
[22]: # Create a list to store the processed labels
all_labels = []

# Iterate over the 'Finding Labels' column
for index, row in df.iterrows():
    labels = row['Finding Labels'].split('|')
    for label in labels:
        all_labels.append(label)

all_labels = list(set(all_labels))
display(f"Number of unique labels: {len(all_labels)}")
```

'Number of unique labels: 7'

```
[23]: # Count occurrences of each label
label_counts = {label: sum(df['Finding Labels'].str.contains(label)) for label_
↪ in all_labels}
```

```

# Convert to DataFrame for plotting
label_counts_df = pd.DataFrame(list(label_counts.items()), columns=['Label', 'Count'])
label_counts_df = label_counts_df.sort_values(by='Count', ascending=False)

# Define a pastel color palette
palette = sns.color_palette("pastel", len(label_counts_df))

# Plot bar chart using pastel colors
plt.figure(figsize=(12, 6))
ax = sns.barplot(x='Label', y='Count', data=label_counts_df, palette=palette)

# Rotate x-axis labels
plt.xticks(rotation=90, fontsize=12)
plt.xlabel('Finding Labels', fontsize=14)
plt.ylabel('Count', fontsize=14)
plt.title('Label Counts in NIH Chest X-rays Dataset', fontsize=16)

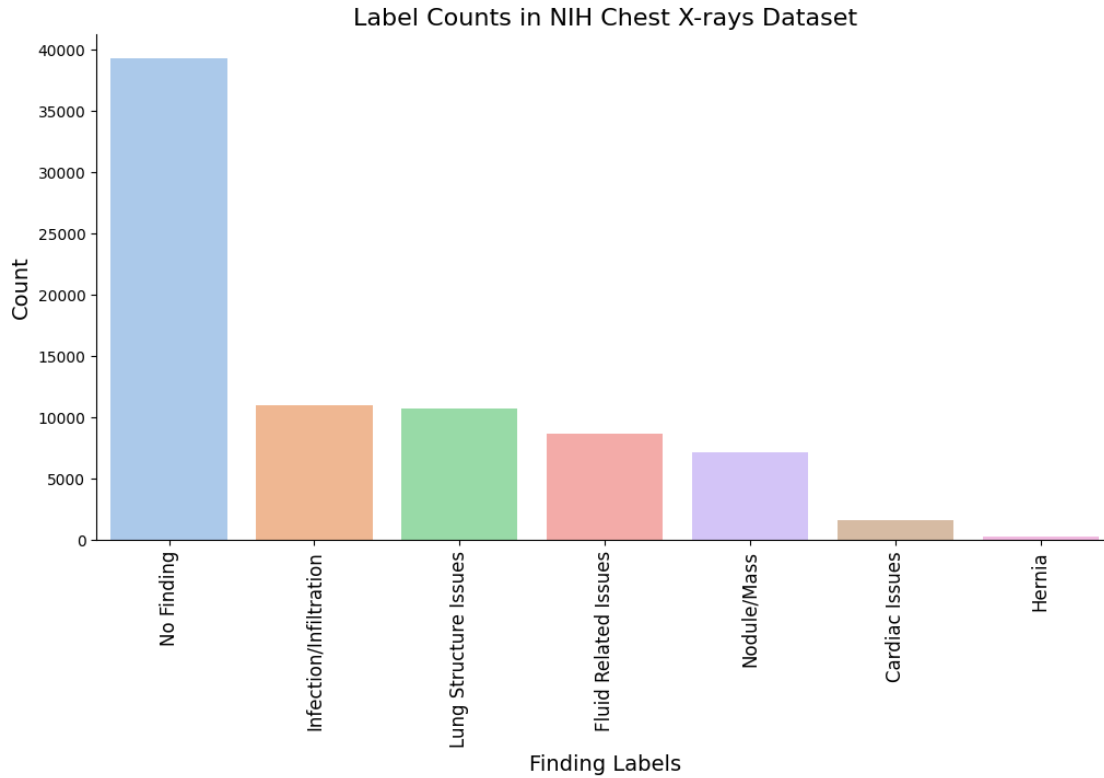
# Remove unnecessary borders
sns.despine()

```

<ipython-input-23-bf94c940165c>:13: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
ax = sns.barplot(x='Label', y='Count', data=label_counts_df, palette=palette)
```



```
[24]: display(label_counts_df)
```

	Label	Count
6	No Finding	39302
1	Infection/Infiltration	10936
2	Lung Structure Issues	10683
5	Fluid Related Issues	8598
4	Nodule/Mass	7140
0	Cardiac Issues	1563
3	Hernia	192

```
[25]: # Count occurrences of each label
label_counts = {label: sum(df['Finding Labels'].str.contains(label)) for label_
↳ in all_labels}

# Convert to DataFrame and sort by count
label_counts_df = pd.DataFrame(list(label_counts.items()), columns=['Label',
↳ 'Count'])
label_counts_df = label_counts_df.sort_values(by='Count', ascending=False)

# Extract the 5 least common labels
least_common_labels_df = label_counts_df.tail(5)
```

```

# Count occurrences for these labels in the dataset
least_common_entries = df[df['Finding Labels']].apply(lambda x: any(label in x_
    ↳for label in least_common_labels_df['Label'])))

# Count the occurrences of these specific labels in the dataset
least_common_label_counts = {label: sum(df['Finding Labels'].str.
    ↳contains(label)) for label in least_common_labels_df['Label']}

# Convert to DataFrame
least_common_label_counts_df = pd.DataFrame(list(least_common_label_counts.
    ↳items()), columns=['Label', 'Count'])
palette = sns.color_palette("pastel")

# Create a bar plot with a custom color scheme
plt.figure(figsize=(10, 5))
ax = sns.barplot(
    x='Label',
    y='Count',
    hue='Label', # Assign hue to avoid warning
    data=least_common_label_counts_df,
    palette=palette,
    legend=False # Remove unnecessary legend
)

# Add count labels above bars
for index, row in least_common_label_counts_df.iterrows():
    ax.text(index, row['Count'] + 50, str(row['Count']), color='black',
    ↳ha="center", fontsize=12)

# Customize plot aesthetics
plt.xticks(rotation=45, fontsize=12)
plt.yticks(fontsize=12)
plt.xlabel('Finding Labels', fontsize=14)
plt.ylabel('Count', fontsize=14)
plt.title('Counts of Least Common Labels in NIH Chest X-rays Dataset',
    ↳fontsize=16)

# Show the plot
plt.tight_layout()
plt.show()

# Remove top and right border for a cleaner look
sns.despine()

# Show the plot
plt.show()

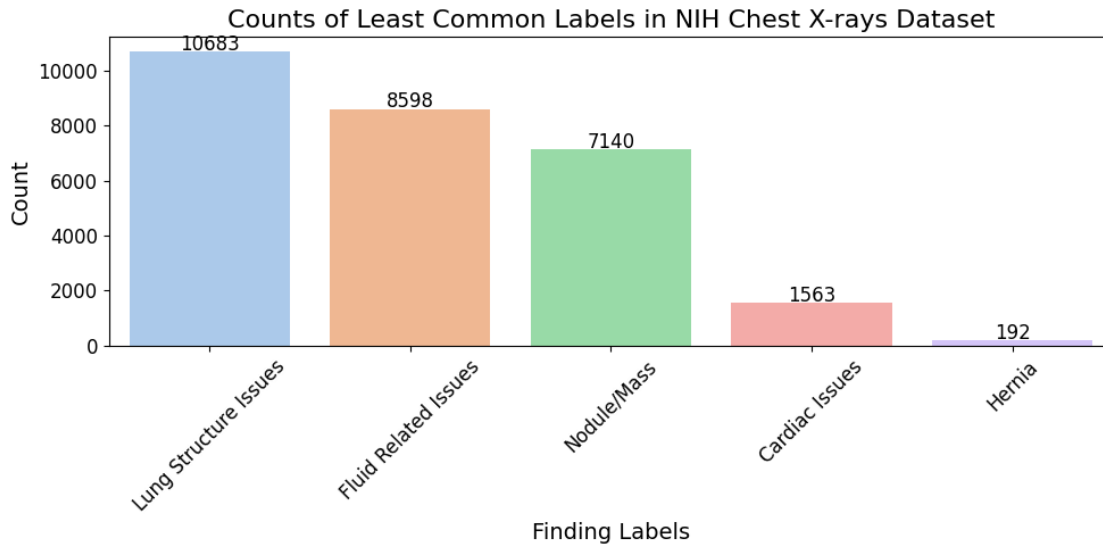
```



```
# Display the number of entries that contain the least common labels
least_common_entries_count = least_common_entries.shape[0]
print(f"Number of entries that contain at least one of the five least common_
↪labels: {least_common_entries_count}")
```

<ipython-input-25-f52b0187c395>:23: UserWarning: The palette list has more values (10) than needed (5), which may not be intended.

```
ax = sns.barplot(
```



<Figure size 640x480 with 0 Axes>

Number of entries that contain at least one of the five least common labels:
22003

0.6 Notes

- The class with lowest number of entries contains 192 images. Since it will be hard to generate synthetic data, reducing all classes to 192 images to address class imbalance will significantly reduce available data.

We can consider the following strategies for addressing the class imbalances.

1. Data Augmentation (For Underrepresented Classes)

Best for Image Data

To increase the number of images in underrepresented classes, we can apply data augmentation techniques such as:

- Geometric Transformations: Rotation, flipping, cropping, translation, scaling.
- Color Variations: Adjust brightness, contrast, saturation.
- Noise Injection: Adding Gaussian noise or slight blur.

- Generative Models: Use GANs (Generative Adversarial Networks) or Diffusion Models to synthesize realistic chest X-ray images.

2. Class Weighting in Model Training

Best for Training Deep Learning Models

Instead of balancing data at the dataset level, we can adjust the model's loss function to give higher weight to underrepresented classes.

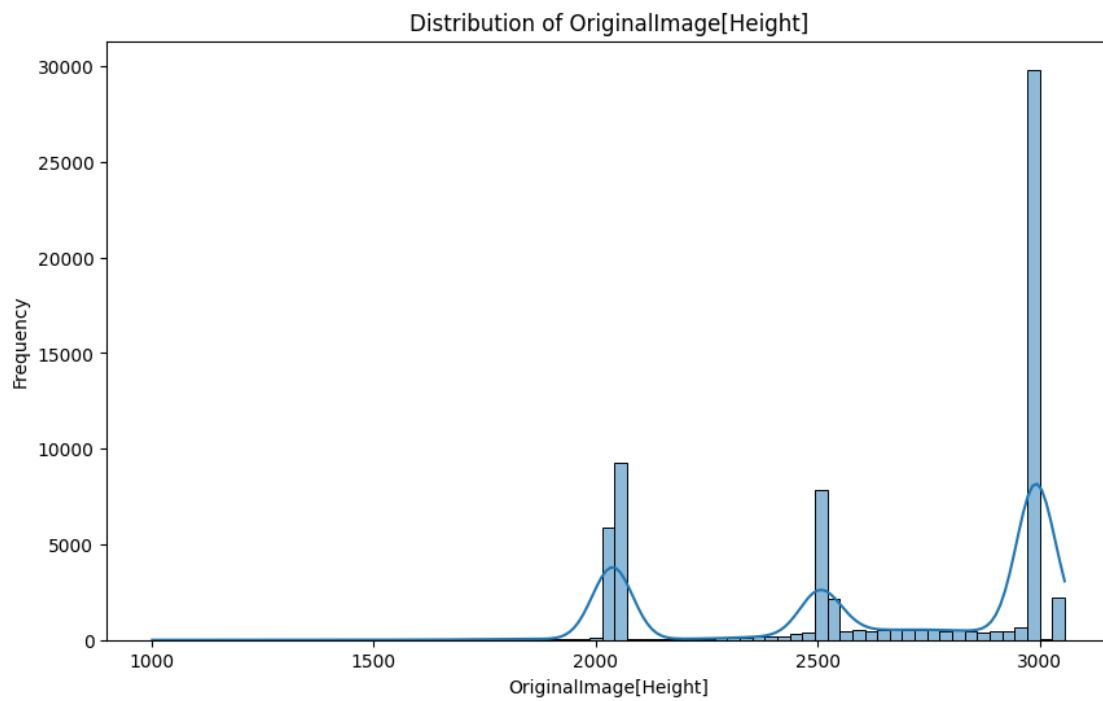
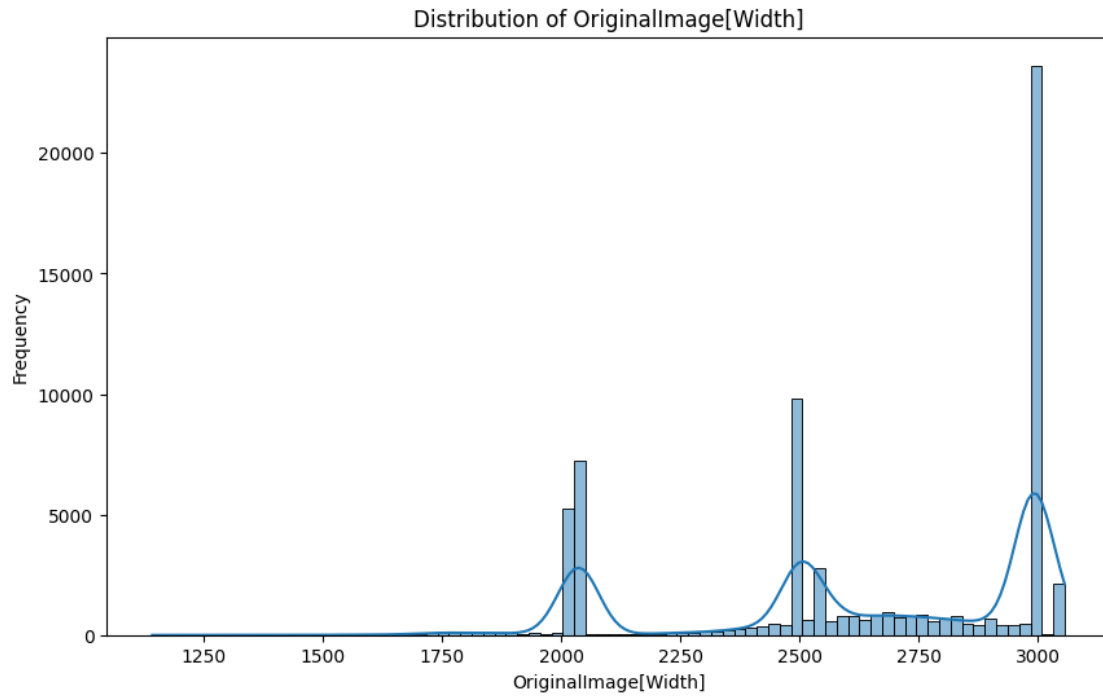
3. Oversampling the Minority Classes

Best when Augmentation is not feasible

Instead of generating new synthetic data, we can duplicate existing samples in the underrepresented classes.

```
[26]: # Plot the distribution of OriginalImage[Width]
plt.figure(figsize=(10, 6))
sns.histplot(df['OriginalImage[Width]'], kde=True)
plt.title('Distribution of OriginalImage[Width]')
plt.xlabel('OriginalImage[Width]')
plt.ylabel('Frequency')
plt.show()

# Plot the distribution of OriginalImage[Height]
plt.figure(figsize=(10, 6))
sns.histplot(df['OriginalImage[Height]'], kde=True)
plt.title('Distribution of OriginalImage[Height]')
plt.xlabel('OriginalImage[Height]')
plt.ylabel('Frequency')
plt.show()
```



0.7 Correlation Matrix

```
[27]: # Create a copy of the DataFrame to avoid modifying the original
df_processed = df.copy()

# Extract each label to a separate boolean column
for label in all_labels:
    df_processed[f'has_{label}'] = df_processed['Finding Labels'].str.
    ↪contains(label)

# drop Finding Labels, Image Index, 'OriginalImage[Width', 'Height]',
    ↪'OriginalImagePixelSpacing[x', 'y]', 'Unnamed: 11'
df_processed = df_processed.drop(columns=['Finding Labels', 'Image Index',
    ↪'OriginalImage[Width', 'Height]', 'OriginalImagePixelSpacing[x', 'y]',
    ↪'Unnamed: 11'])

display(df_processed.head())
# Encode categorical columns using one-hot encoding
categorical_columns = ['Patient Gender', 'View Position']
for column in categorical_columns:
    df_processed = pd.get_dummies(df_processed, columns=[column], prefix=[column])

# Encode the label for the boolean columns
for label in all_labels:
    df_processed[f'has_{label}'] = df_processed[f'has_{label}'].astype(int)

# Build the correlation matrix
correlation_matrix = df_processed.corr()

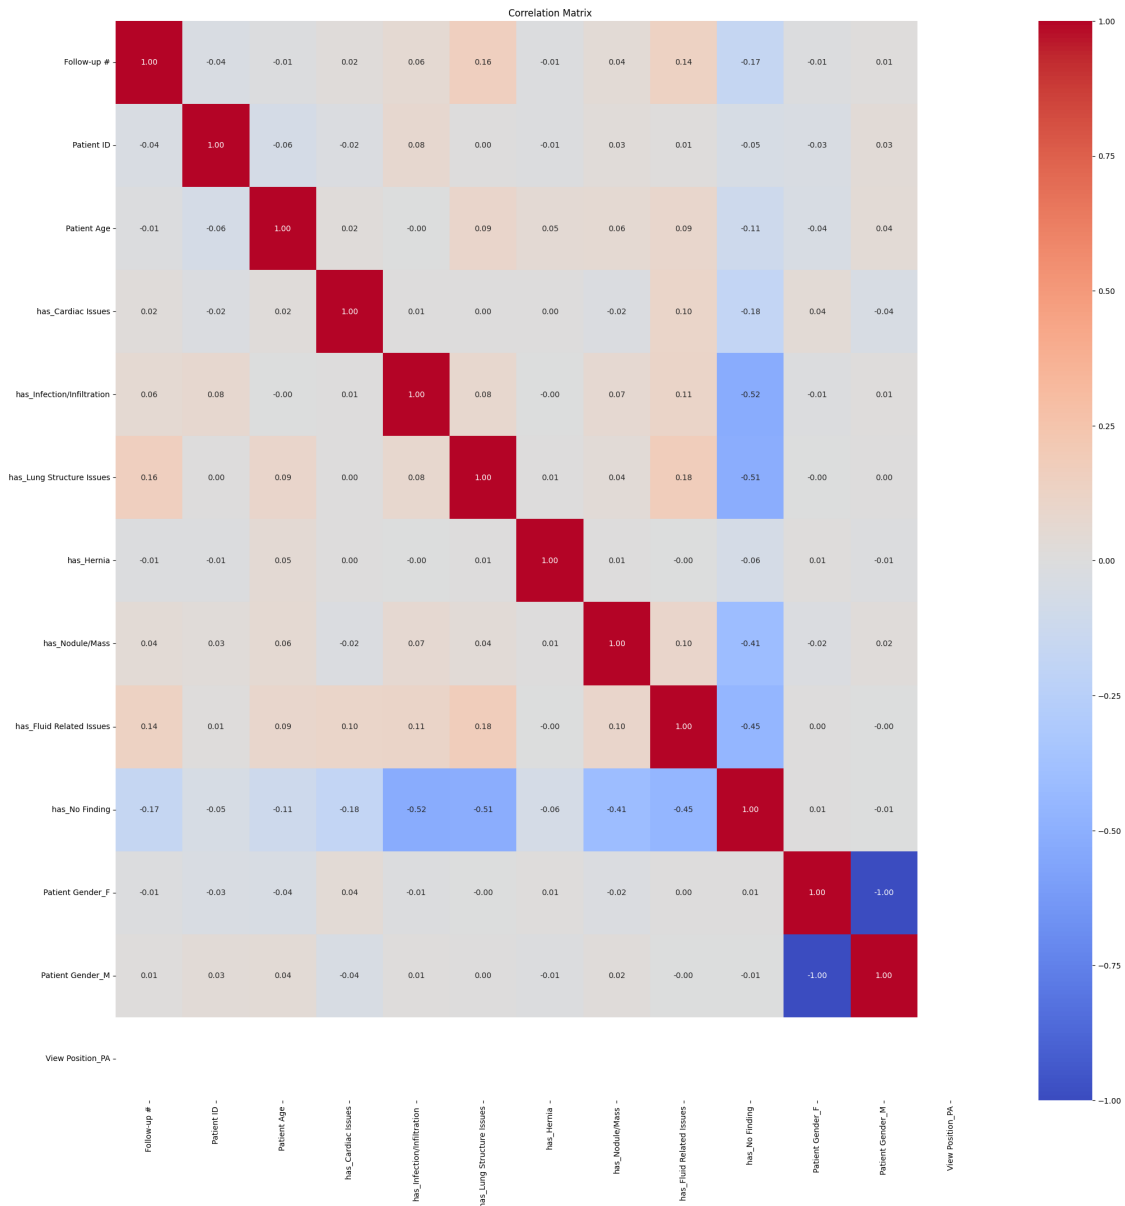
# Plot correlation matrix
plt.figure(figsize=(25, 25))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix')
plt.show()
```

	Follow-up #	Patient ID	Patient Age	Patient Gender	View Position	\
0	0	1	58	M	PA	
1	1	1	58	M	PA	
2	2	1	58	M	PA	
3	0	2	81	M	PA	
4	0	3	81	F	PA	

	has_Cardiac Issues	has_Infection/Infiltration	has_Lung Structure Issues	\
0	True	False	False	
1	True	False	True	
2	True	False	False	

3	False	False	False
4	False	False	False

	has_Hernia	has_Nodule/Mass	has_Fluid Related Issues	has_No Finding
0	False	False	False	False
1	False	False	False	False
2	False	False	True	False
3	False	False	False	True
4	True	False	False	False



```
[28]: display(correlation_matrix)
```

	Follow-up #	Patient ID	Patient Age \
Follow-up #	1.000000	-0.037928	-0.013289
Patient ID	-0.037928	1.000000	-0.063590
Patient Age	-0.013289	-0.063590	1.000000
has_Cardiac Issues	0.018726	-0.019187	0.015911
has_Infection/Infiltration	0.058027	0.077087	-0.001701
has_Lung Structure Issues	0.161326	0.004941	0.094682
has_Hernia	-0.009183	-0.013524	0.051530
has_Nodule/Mass	0.044365	0.026500	0.055799
has_Fluid Related Issues	0.139019	0.013117	0.087079
has_No Finding	-0.166038	-0.051644	-0.112697
Patient Gender_F	-0.009111	-0.033435	-0.044918
Patient Gender_M	0.009111	0.033435	0.044918
View Position_PA	NaN	NaN	NaN

	has_Cardiac Issues	has_Infection/Infiltration \
Follow-up #	0.018726	0.058027
Patient ID	-0.019187	0.077087
Patient Age	0.015911	-0.001701
has_Cardiac Issues	1.000000	0.007503
has_Infection/Infiltration	0.007503	1.000000
has_Lung Structure Issues	0.002681	0.079501
has_Hernia	0.002851	-0.000902
has_Nodule/Mass	-0.020437	0.067088
has_Fluid Related Issues	0.104129	0.110167
has_No Finding	-0.182645	-0.521742
Patient Gender_F	0.039572	-0.012901
Patient Gender_M	-0.039572	0.012901
View Position_PA	NaN	NaN

	has_Lung Structure Issues	has_Hernia \
Follow-up #	0.161326	-0.009183
Patient ID	0.004941	-0.013524
Patient Age	0.094682	0.051530
has_Cardiac Issues	0.002681	0.002851
has_Infection/Infiltration	0.079501	-0.000902
has_Lung Structure Issues	1.000000	0.006501
has_Hernia	0.006501	1.000000
has_Nodule/Mass	0.038921	0.007810
has_Fluid Related Issues	0.177758	-0.001273
has_No Finding	-0.514519	-0.063357
Patient Gender_F	-0.002363	0.014536
Patient Gender_M	0.002363	-0.014536
View Position_PA	NaN	NaN

	has_Nodule/Mass	has_Fluid Related Issues \
Follow-up #	0.044365	0.139019
Patient ID	0.026500	0.013117

Patient Age	0.055799	0.087079
has_Cardiac Issues	-0.020437	0.104129
has_Infection/Infiltration	0.067088	0.110167
has_Lung Structure Issues	0.038921	0.177758
has_Hernia	0.007810	-0.001273
has_Nodule/Mass	1.000000	0.100158
has_Fluid Related Issues	0.100158	1.000000
has_No Finding	-0.408061	-0.453317
Patient Gender_F	-0.019791	0.000407
Patient Gender_M	0.019791	-0.000407
View Position_PA	NaN	NaN

	has_No Finding	Patient Gender_F \
Follow-up #	-0.166038	-0.009111
Patient ID	-0.051644	-0.033435
Patient Age	-0.112697	-0.044918
has_Cardiac Issues	-0.182645	0.039572
has_Infection/Infiltration	-0.521742	-0.012901
has_Lung Structure Issues	-0.514519	-0.002363
has_Hernia	-0.063357	0.014536
has_Nodule/Mass	-0.408061	-0.019791
has_Fluid Related Issues	-0.453317	0.000407
has_No Finding	1.000000	0.005735
Patient Gender_F	0.005735	1.000000
Patient Gender_M	-0.005735	-1.000000
View Position_PA	NaN	NaN

	Patient Gender_M	View Position_PA
Follow-up #	0.009111	NaN
Patient ID	0.033435	NaN
Patient Age	0.044918	NaN
has_Cardiac Issues	-0.039572	NaN
has_Infection/Infiltration	0.012901	NaN
has_Lung Structure Issues	0.002363	NaN
has_Hernia	-0.014536	NaN
has_Nodule/Mass	0.019791	NaN
has_Fluid Related Issues	-0.000407	NaN
has_No Finding	-0.005735	NaN
Patient Gender_F	-1.000000	NaN
Patient Gender_M	1.000000	NaN
View Position_PA	NaN	NaN

0.8 Additional Charts and EDA

```
[29]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

```

# Create a copy of the DataFrame to avoid modifying the original
df_processed = df.copy()

# Extract each label to a separate boolean column
for label in all_labels:
    df_processed[f'has_{label}'] = df_processed['Finding Labels'].fillna('').
    ↪str.contains(label, regex=False).astype(int)

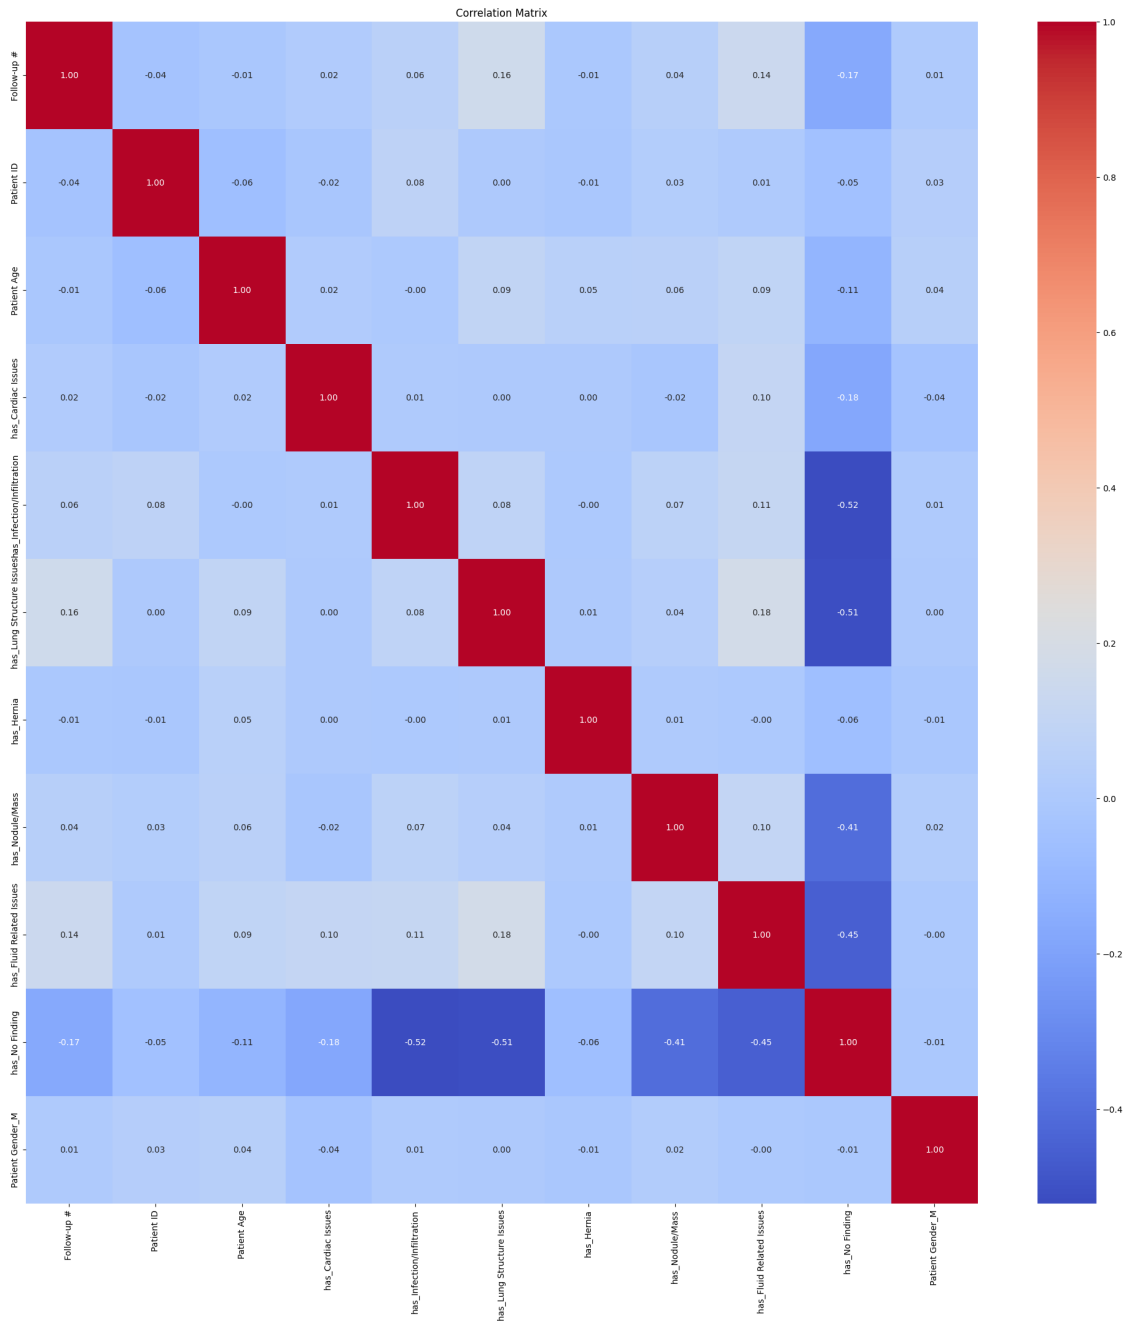
# Drop unnecessary columns
drop_columns = ['Finding Labels', 'Image Index', 'OriginalImage[Width',
    ↪'Height]',
    'OriginalImagePixelSpacing[x', 'y]', 'Unnamed: 11']
df_processed = df_processed.drop(columns=drop_columns, errors='ignore')

# Encode categorical columns using one-hot encoding, dropping first category to
    ↪avoid redundancy
categorical_columns = ['Patient Gender', 'View Position']
df_processed = pd.get_dummies(df_processed, columns=categorical_columns,
    ↪drop_first=True)

# Build the correlation matrix
correlation_matrix = df_processed.corr()

# Plot correlation matrix
plt.figure(figsize=(25, 25))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")
plt.title('Correlation Matrix')
plt.show()

```

```
[30]: import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

# Extract disease label columns, excluding 'No Finding'
disease_columns = [col for col in df_processed.columns if col.
    ↳startswith('has_') and col != 'has_No Finding']
```

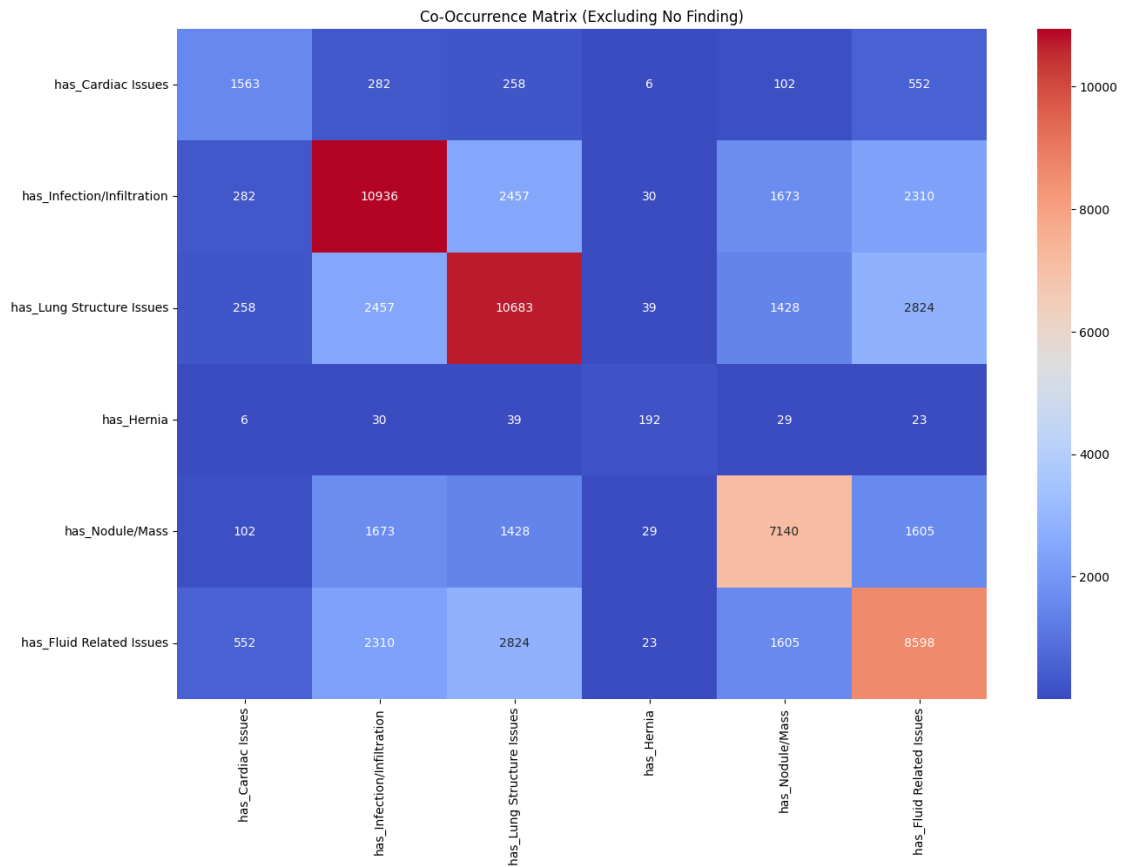
```

# Compute the co-occurrence matrix for the remaining disease labels
co_occurrence_matrix = np.dot(df_processed[disease_columns].T,
    ↪df_processed[disease_columns])

# Convert to a DataFrame for easier visualization
co_occurrence_df = pd.DataFrame(co_occurrence_matrix, index=disease_columns,
    ↪columns=disease_columns)

# Plot the co-occurrence heatmap
plt.figure(figsize=(15, 10))
sns.heatmap(co_occurrence_df, annot=True, fmt=".0f", cmap="coolwarm")
plt.title('Co-Occurrence Matrix (Excluding No Finding)')
plt.show()

```



0.8.1 Preview Annotated Image

```
[31]: first_image = df_box_list.iloc[0]
image_info = df_box_list[df_box_list['Image Index'] == first_image['Image_
↳Index']]

display(image_info)

# Define image file name
image_name = first_image['Image Index']

# Determine the folder using dictionary lookup
image_folder = get_image_folder(df_locked, image_name)
folder_number = int(image_folder.split("_")[1])

display(f"Image Name: {image_name}")
display(image_folder)
display(folder_number)

if SKIP_BOUNDING_BOX == False:
    urllib.request.urlretrieve(links[folder_number - 1], f"{image_folder}.tar.
↳gz") # download the zip file
```

```
      Image Index Finding Label          x          y          w \
0  00013118_008.png  Atelectasis  225.084746  547.019217  86.779661

      h  Unnamed: 6  Unnamed: 7  Unnamed: 8
0  79.186441      NaN          NaN          NaN

'Image Name: 00013118_008.png'
'images_006'
6
```

```
[32]: if SKIP_BOUNDING_BOX == False:
      !tar -xvzf images_006.tar.gz
```

```
[33]: # Construct the image file path
image_file = f"images/{image_name}"

if SKIP_BOUNDING_BOX == False:
    try:
        image = cv2.imread(image_file)
        if image is None:
            print(f"Error: Could not read image file '{image_file}'")
        else:
            image_info = df_box_list[df_box_list['Image Index'] == image_name]
            display(image_info.columns)
```

```

    # Extract bounding box values as integers
    x = int(image_info['x'].iloc[0])
    y = int(image_info['y'].iloc[0])
    w = int(image_info['w'].iloc[0])
    h = int(image_info['h'].iloc[0])

    # Draw the bounding box on the image
    cv2.rectangle(image, (x, y), (x + w, y + h), (255, 0, 0), 4)

    # Convert BGR to RGB for displaying
    plt.figure(figsize=(10, 10))
    plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
    plt.title(f"Bounding Box on {image_name}")
    plt.axis("off")
    plt.show()
except Exception as e:
    print(f"An error occurred: {e}")

```

0.9 Pre-processing - Image Scaling

```

[34]: # Reduce rows with 'No Finding' label to a maximum of 10,000
no_finding_df = df[df['Finding Labels'] == 'No Finding']
if len(no_finding_df) > 10000:
    no_finding_df = no_finding_df.sample(n=10000, random_state=42) # Randomly
    ↪sample 10,000 rows

# Concatenate the reduced 'No Finding' rows with other rows
other_findings_df = df[df['Finding Labels'] != 'No Finding']
df = pd.concat([no_finding_df, other_findings_df], ignore_index=True)

```

```

[35]: len(df)

```

```

[35]: 38008

```

```

[36]: # Rename columns
df = df.rename(columns={
    "OriginalImage[Width]": "width",
    "Height": "height",
    "OriginalImagePixelSpacing[x]": "pixel_spacing x",
    "y": "pixel_spacing y"
})

display(df.head())

```

	Image Index	Finding Labels	Follow-up #	Patient ID	Patient Age \
0	00019856_000.png	No Finding	0	19856	57
1	00001020_000.png	No Finding	0	1020	52

2	00008187_001.png	No Finding	1	8187	59
3	00003360_003.png	No Finding	3	3360	8
4	00014364_000.png	No Finding	0	14364	26

	Patient	Gender	View	Position	width	height	pixel_spacing	x \
0		M		PA	2992	2991	0.143	
1		M		PA	2500	2048	0.171	
2		M		PA	2500	2048	0.168	
3		M		PA	2048	2500	0.168	
4		F		PA	2454	2991	0.143	

	pixel_spacing	y	Unnamed: 11
0	0.143		NaN
1	0.171		NaN
2	0.168		NaN
3	0.168		NaN
4	0.143		NaN

```
[37]: # drop 'Unnamed: 11' column
df = df.drop(columns=['Unnamed: 11'], errors='ignore')
display(df.head())
```

	Image	Index	Finding	Labels	Follow-up	#	Patient	ID	Patient	Age \
0	00019856_000.png		No Finding		0		19856			57
1	00001020_000.png		No Finding		0		1020			52
2	00008187_001.png		No Finding		1		8187			59
3	00003360_003.png		No Finding		3		3360			8
4	00014364_000.png		No Finding		0		14364			26

	Patient	Gender	View	Position	width	height	pixel_spacing	x \
0		M		PA	2992	2991	0.143	
1		M		PA	2500	2048	0.171	
2		M		PA	2500	2048	0.168	
3		M		PA	2048	2500	0.168	
4		F		PA	2454	2991	0.143	

	pixel_spacing	y
0	0.143	
1	0.171	
2	0.168	
3	0.168	
4	0.143	

```
[38]: import os
import tarfile
import urllib.request

def download_and_extract(links, folder_ranges, df_locked):
```

```

"""Downloads image archives, extracts them, and organizes images."""

if not os.path.exists("images"):
    os.makedirs("images")

for i, link in enumerate(links):
    folder_name = f"images_{i+1:03d}"
    archive_name = f"{folder_name}.tar.gz"

    if not os.path.exists(os.path.join("images", archive_name)): #check if
↳ the archive already exists to prevent unnecessary downloads
        print(f"Downloading {archive_name}...")
        urllib.request.urlretrieve(link, archive_name)
    else:
        print(f"Skipping download for {archive_name} as file already exists")

    try:
        print(f"Extracting {archive_name}...")
        with tarfile.open(archive_name, "r:gz") as tar:
            tar.extractall()
        print("Extraction complete.")
    except Exception as e:
        print(f"Error extracting {archive_name}: {e}")
        continue # Skip to the next archive if extraction fails

# Move extracted images to the 'images' folder
source_folder = folder_name
if os.path.exists(source_folder):
    extracted_files = os.listdir(source_folder)
    for file in extracted_files:
        source_file = os.path.join(source_folder, file)
        destination_file = os.path.join("images", file)
        try:
            os.rename(source_file, destination_file)
        except FileExistsError:
            print(f"File {file} already exists in images folder, skipping")

    os.rmdir(source_folder)
else:
    print(f"Folder {source_folder} doesn't exist")

# Remove the archive file
try:
    os.remove(archive_name)
    print(f"Removed {archive_name}")

```

```

        except OSError as e:
            print(f"Error removing {archive_name}: {e}")

if SKIP_DOWNLOAD == False:
    download_and_extract(links, folder_ranges, df_locked)

```

```

[39]: if SKIP_DOWNLOAD == False:
        image_folder = 'images'
        num_images = len([f for f in os.listdir(image_folder) if os.path.isfile(os.
↪path.join(image_folder, f))])
        print(f"Number of images in '{image_folder}' folder: {num_images}")

```

```

[40]: if SKIP_DOWNLOAD == False:
        image_folder = 'images'

        # Get a set of image names from the 'Image Index' column of the DataFrame
        image_names_in_df = set(df['Image Index'].unique())

        print(len(image_names_in_df))

        # Iterate through all files in the image folder
        for filename in os.listdir(image_folder):
            filepath = os.path.join(image_folder, filename)

            # Check if it's a file and not in the DataFrame's 'Image Index' column
            if os.path.isfile(filepath) and filename not in image_names_in_df:
                try:
                    os.remove(filepath)
                    print(f"Removed file: {filename}")
                except OSError as e:
                    print(f"Error deleting file {filename}: {e}")

```

```

[41]: image_folder = 'images'
def get_num_images(image_folder):
    num_images = len([f for f in os.listdir(image_folder) if os.path.isfile(os.
↪path.join(image_folder, f))])
    return num_images

if SKIP_DOWNLOAD == False:
    print(f"Number of images in '{image_folder}' folder:
↪{get_num_images(image_folder)}")

```

```

[42]: if SKIP_DOWNLOAD == False:
        !python image_scale.py

```

```

[43]: if SKIP_DOWNLOAD == False:
        !zip -r images_resized.zip images_resized

```

```
[44]: if SKIP_DOWNLOAD == False:
    print(f"Number of images in 'images_resized' folder:␣
    ↪{get_num_images('images_resized')}")
```

```
[45]: def zip_folder(folder_path, zip_filename):
    """Zips a folder.
    Args:
        folder_path: The path to the folder to zip.
        zip_filename: The name of the zip file to create.
    """

    # Create a zip archive
    with zipfile.ZipFile(zip_filename, 'w', zipfile.ZIP_DEFLATED) as zipf:
        for root, _, files in os.walk(folder_path):
            for file in files:
                zipf.write(os.path.join(root, file),
                           os.path.relpath(os.path.join(root, file),
                                               os.path.join(folder_path, '..')))

    if SKIP_DOWNLOAD == False:
        zip_folder('images_resized', 'images_resized.zip')
```

```
[46]: if SKIP_DOWNLOAD == False:
    !cp images_resized.zip {RESIZED_IMAGES_ZIP_PATH}
```

```
[47]: import zipfile
import os

def unzip_files(zip_path, extract_path):
    """Unzips files from a zip archive to a specified directory.

    Args:
        zip_path: Path to the zip file.
        extract_path: Directory to extract the files to.
    """
    try:
        with zipfile.ZipFile(zip_path, 'r') as zip_ref:
            zip_ref.extractall(extract_path)
        print(f"Successfully unzipped '{zip_path}' to '{extract_path}'")
    except FileNotFoundError:
        print(f"Error: Zip file not found at '{zip_path}'")
    except zipfile.BadZipFile:
        print(f"Error: Invalid zip file at '{zip_path}'")
    except Exception as e:
        print(f"An unexpected error occurred: {e}")

# Assuming RESIZED_IMAGES_ZIP_PATH is defined and holds the correct path
if SKIP_UNZIP == False:
```



```
unzip_files(REIZED_IMAGES_ZIP_PATH + "/images_resized.zip", "images_resized")
```

Successfully unzipped

```
'/content/drive/MyDrive/AAI-590_Collabs/images_resized.zip' to 'images_resized'
```

0.9.1 Multi-label encoding

```
[48]: # Extract all unique labels
all_labels = sorted(set(label for sublist in df['Finding Labels'].str.
    ↪split('|') for label in sublist))
display(all_labels)

# Encode multi-labels
def encode_multilabel(labels):
    label_set = labels.split('|')
    return [1 if label in label_set else 0 for label in all_labels]

df['encoded_labels'] = df['Finding Labels'].apply(encode_multilabel)
display(df.head())

y = np.array(df['encoded_labels'].tolist())
```

```
['Cardiac Issues',
 'Fluid Related Issues',
 'Hernia',
 'Infection/Infiltration',
 'Lung Structure Issues',
 'No Finding',
 'Nodule/Mass']
```

	Image Index	Finding Labels	Follow-up #	Patient ID	Patient Age	\
0	00019856_000.png	No Finding	0	19856	57	
1	00001020_000.png	No Finding	0	1020	52	
2	00008187_001.png	No Finding	1	8187	59	
3	00003360_003.png	No Finding	3	3360	8	
4	00014364_000.png	No Finding	0	14364	26	

	Patient Gender	View Position	width	height	pixel_spacing x	\
0	M	PA	2992	2991	0.143	
1	M	PA	2500	2048	0.171	
2	M	PA	2500	2048	0.168	
3	M	PA	2048	2500	0.168	
4	F	PA	2454	2991	0.143	

	pixel_spacing y	encoded_labels
0	0.143	[0, 0, 0, 0, 0, 1, 0]
1	0.171	[0, 0, 0, 0, 0, 1, 0]
2	0.168	[0, 0, 0, 0, 0, 1, 0]

```
3          0.168 [0, 0, 0, 0, 0, 1, 0]
4          0.143 [0, 0, 0, 0, 0, 1, 0]
```

0.9.2 Encode Tabular Labels

```
[49]: # Encode gender (e.g., Male/Female -> 0/1)
df['Patient Gender'] = LabelEncoder().fit_transform(df['Patient Gender'])

# Standardize numerical features
scaler = StandardScaler()
df['Patient Age'] = scaler.fit_transform(df[['Patient Age']])
df['Follow-up #'] = scaler.fit_transform(df[['Follow-up #']])

display(df.head())
```

	Image Index	Finding Labels	Follow-up #	Patient ID	Patient Age \
0	00019856_000.png	No Finding	-0.552742	19856	0.525833
1	00001020_000.png	No Finding	-0.552742	1020	0.215450
2	00008187_001.png	No Finding	-0.457542	8187	0.649986
3	00003360_003.png	No Finding	-0.267142	3360	-2.515918
4	00014364_000.png	No Finding	-0.552742	14364	-1.398540

	Patient Gender	View Position	width	height	pixel_spacing x \
0	1	PA	2992	2991	0.143
1	1	PA	2500	2048	0.171
2	1	PA	2500	2048	0.168
3	1	PA	2048	2500	0.168
4	0	PA	2454	2991	0.143

	pixel_spacing y	encoded_labels
0	0.143	[0, 0, 0, 0, 0, 1, 0]
1	0.171	[0, 0, 0, 0, 0, 1, 0]
2	0.168	[0, 0, 0, 0, 0, 1, 0]
3	0.168	[0, 0, 0, 0, 0, 1, 0]
4	0.143	[0, 0, 0, 0, 0, 1, 0]

Train/Test Split

```
[50]: # Create a copy of the DataFrame
df_processed = df.copy()

# Drop unnecessary columns
columns_to_drop = ['Finding Labels', 'width', 'height', 'pixel_spacing x',
                  'pixel_spacing y', 'View Position']
df_processed = df_processed.drop(columns=columns_to_drop, errors='ignore')

# Perform train/validation split
```

```

train_df, val_df = train_test_split(df_processed, test_size=0.2,
    ↪random_state=42)
print("Train size:", len(train_df))
print("Val size:", len(val_df))

display(train_df.head())

```

Train size: 30406

Val size: 7602

	Image Index	Follow-up #	Patient ID	Patient Age	Patient Gender	\
2855	00015255_013.png	0.684857	15255	-0.157009		1
6626	00008663_002.png	-0.362342	8663	-0.467392		0
12335	00002471_001.png	-0.457542	2471	-0.157009		1
11355	00001397_002.png	-0.362342	1397	0.339603		1
30189	00020495_000.png	-0.552742	20495	0.153374		0

	encoded_labels
2855	[0, 0, 0, 0, 0, 1, 0]
6626	[0, 0, 0, 0, 0, 1, 0]
12335	[0, 1, 0, 1, 1, 0, 0]
11355	[0, 1, 0, 0, 0, 0, 0]
30189	[0, 0, 0, 1, 0, 0, 0]

0.9.3 Address Class Imbalance

```

[51]: # Flatten all labels into one array for each class
class_weights = {}
for i, label in enumerate(all_labels):
    label_values = y[:, i]
    class_weights[i] = compute_class_weight(class_weight='balanced', classes=np.
    ↪array([0, 1]), y=label_values)[1] # 1 = positive class, classes is now a
    ↪NumPy array

print(class_weights)

```

```

{0: np.float64(12.15866922584773), 1: np.float64(2.2102814608048384), 2:
np.float64(98.97916666666667), 3: np.float64(1.7377468910021945), 4:
np.float64(1.7789010577553122), 5: np.float64(1.9004), 6:
np.float64(2.661624649859944)}

```

0.9.4 Data Generator (Image + Tabular)

```

[52]: import tensorflow as tf
import os
import numpy as np

IMG_SIZE = 512 # Reduced image size

```

```

IMAGE_PATH = '/content/images_resized/images_resized'

def load_image(image_id):
    """Loads and preprocesses a single image."""
    # Convert image_id to string before joining paths
    path = os.path.join(IMAGE_PATH, image_id.numpy().decode('utf-8')) # Decodes
    ↪ the byte string to a regular string
    image = tf.io.read_file(path)
    image = tf.image.decode_jpeg(image, channels=3)
    image = tf.image.resize(image, [IMG_SIZE, IMG_SIZE]) # Resize the image to
    ↪ a fixed size
    image = image / 255.0 # Normalize pixel values
    return image

def make_dataset(df):
    """Creates a tf.data.Dataset for the given DataFrame."""
    image_paths = df['Image Index'].values
    tabular_features = df[['Follow-up #', 'Patient Age', 'Patient Gender']].
    ↪ values.astype(np.float32)
    labels = np.array(df['encoded_labels'].tolist()).astype(np.float32)

    # Create a tf.data.Dataset from image paths and labels
    dataset = tf.data.Dataset.from_tensor_slices((image_paths,
    ↪ tabular_features, labels))
    # Map the load_image function to load and preprocess images on the fly
    dataset = dataset.map(lambda img_path, tab_feat, label: (tf.
    ↪ py_function(load_image, [img_path], tf.float32), tab_feat, label), # Use tf.
    ↪ py_function
                        num_parallel_calls=tf.data.AUTOTUNE) # Use parallel
    ↪ calls for faster loading

    # Set the shape of the image tensor explicitly
    dataset = dataset.map(lambda image, tab_feat, label: (tf.
    ↪ ensure_shape(image, [IMG_SIZE, IMG_SIZE, 3]), tab_feat, label))

    # Batch and prefetch the dataset
    dataset = dataset.batch(4).prefetch(tf.data.AUTOTUNE)
    return dataset

train_dataset = make_dataset(train_df)
val_dataset = make_dataset(val_df)

```

0.9.5 Build the Hybrid CNN + Tabular Model

```
[53]: def build_model(hp, img_shape=(512, 512, 3), tab_shape=(3,)),  
      ↪ num_labels=len(all_labels)):  
      # Image branch  
      img_input = Input(shape=img_shape, name='input_layer')  
      x = layers.Conv2D(hp.Int('conv_1_filters', min_value=32, max_value=128,  
      ↪ step=32), (3, 3), activation='relu')(img_input)  
      x = layers.MaxPooling2D()(x)  
      x = layers.Conv2D(hp.Int('conv_2_filters', min_value=64, max_value=256,  
      ↪ step=64), (3, 3), activation='relu')(x)  
      x = layers.MaxPooling2D()(x)  
      x = layers.Flatten()(x)  
  
      # Tabular branch  
      tab_input = Input(shape=tab_shape, name='input_layer_1')  
      t = layers.Dense(hp.Int('dense_units', min_value=32, max_value=128,  
      ↪ step=32), activation='relu')(tab_input)  
  
      # Combine  
      combined = layers.concatenate([x, t])  
      z = layers.Dense(hp.Int('dense_units_2', min_value=64, max_value=256,  
      ↪ step=64), activation='relu')(combined)  
      z = layers.Dropout(hp.Float('dropout_rate', min_value=0.2, max_value=0.5,  
      ↪ step=0.1))(z)  
      output = layers.Dense(num_labels, activation='sigmoid')(z)  
  
      model = Model(inputs=[img_input, tab_input], outputs=output)  
      model.compile(optimizer='adam', loss='binary_crossentropy',  
      ↪ metrics=['binary_accuracy'])  
      return model
```

0.9.6 Train the Model

```
[54]: # Custom training loop  
def format_batch(batch):  
    # batch is now a tuple of (image, tab, label)  
    image, tab, label = batch  
    return ({"input_layer": image, "input_layer_1": tab}, label)  
  
train_ds = train_dataset.map(lambda x,y,z: format_batch((x,y,z))) # Pass all 3  
    ↪ elements as a single tuple to format_batch  
val_ds = val_dataset.map(lambda x,y,z: format_batch((x,y,z))) # Pass all 3  
    ↪ elements as a single tuple to format_batch
```

```
[55]: DRIVE_TUNER_RESULTS_DIR = DRIVE_PATH + '/tuner_results'
```

```
# Create the directory if it doesn't exist
!mkdir -p "{DRIVE_TUNER_RESULTS_DIR}"
```

```
[56]: # Create a HyperModel instance
class CustomHyperModel(HyperModel):
    def build(self, hp):
        return build_model(hp)

# Set up Keras Tuner Random Search
tuner = RandomSearch(
    CustomHyperModel(),
    objective='val_binary_accuracy', # Optimize for validation binary accuracy
    max_trials=10, # Number of random search trials
    executions_per_trial=3, # Number of executions per trial
    directory=DRIVE_TUNER_RESULTS_DIR, # Directory to save results
    project_name='image_and_tabular_tuning'
)

# Train the tuner to search for the best hyperparameters
tuner.search(train_ds, validation_data=val_ds, epochs=3,
    ↪class_weight=class_weights)

# Get the best model from the search
best_model = tuner.get_best_models(num_models=1)[0]

# Display the best model's summary
best_model.summary()
```

Trial 10 Complete [00h 37m 09s]

Best val_binary_accuracy So Far: 0.8146915833155314

Total elapsed time: 1d 01h 23m 50s

/usr/local/lib/python3.11/dist-packages/keras/src/saving/saving_lib.py:757:

UserWarning: Skipping variable loading for optimizer 'adam', because it has 2 variables whereas the saved optimizer has 22 variables.

saveable.load_own_variables(weights_store.get(inner_path))

Model: "functional"

Layer (type)	Output Shape	Param #	Connected
↪to			
input_layer (InputLayer)	(None, 512, 512, 3)	0	-
↪			

conv2d (Conv2D)	(None, 510, 510, 128)	3,584	┐
↳input_layer[0][0]			
max_pooling2d	(None, 255, 255, 128)	0	┐
↳conv2d[0][0]			
(MaxPooling2D)			┐
↳			
conv2d_1 (Conv2D)	(None, 253, 253, 128)	147,584	┐
↳max_pooling2d[0][0]			
max_pooling2d_1	(None, 126, 126, 128)	0	┐
↳conv2d_1[0][0]			
(MaxPooling2D)			┐
↳			
input_layer_1	(None, 3)	0	- ┐
↳			
(InputLayer)			┐
↳			
flatten (Flatten)	(None, 2032128)	0	┐
↳max_pooling2d_1[0][0]			
dense (Dense)	(None, 32)	128	┐
↳input_layer_1[0][0]			
concatenate (Concatenate)	(None, 2032160)	0	┐
↳flatten[0][0],			
			┐
↳dense[0][0]			
dense_1 (Dense)	(None, 128)	260,116,608	┐
↳concatenate[0][0]			
dropout (Dropout)	(None, 128)	0	┐
↳dense_1[0][0]			
dense_2 (Dense)	(None, 7)	903	┐
↳dropout[0][0]			

Total params: 260,268,807 (992.85 MB)

Trainable params: 260,268,807 (992.85 MB)

Non-trainable params: 0 (0.00 B)

```
[58]: # Continue training the best model for 10 epochs
best_model.fit(train_ds, validation_data=val_ds, epochs=10,
               class_weight=class_weights)
```

Epoch 1/10

7602/7602 326s 43ms/step

- binary_accuracy: 0.8148 - loss: 1.3812 - val_binary_accuracy: 0.8144 -
val_loss: 0.4553

Epoch 2/10

7602/7602 300s 39ms/step

- binary_accuracy: 0.8146 - loss: 1.3894 - val_binary_accuracy: 0.8147 -
val_loss: 0.4565

Epoch 3/10

7602/7602 308s 41ms/step

- binary_accuracy: 0.8151 - loss: 1.3890 - val_binary_accuracy: 0.8146 -
val_loss: 0.4541

Epoch 4/10

7602/7602 299s 39ms/step

- binary_accuracy: 0.8150 - loss: 1.3863 - val_binary_accuracy: 0.8146 -
val_loss: 0.4535

Epoch 5/10

7602/7602 309s 41ms/step

- binary_accuracy: 0.8150 - loss: 1.3928 - val_binary_accuracy: 0.8146 -
val_loss: 0.4507

Epoch 6/10

7602/7602 296s 39ms/step

- binary_accuracy: 0.8149 - loss: 1.3933 - val_binary_accuracy: 0.8143 -
val_loss: 0.4538

Epoch 7/10

7602/7602 308s 41ms/step

- binary_accuracy: 0.8148 - loss: 1.3857 - val_binary_accuracy: 0.8145 -
val_loss: 0.4505

Epoch 8/10

7602/7602 297s 39ms/step

- binary_accuracy: 0.8146 - loss: 1.3919 - val_binary_accuracy: 0.8148 -
val_loss: 0.4506

Epoch 9/10

7602/7602 310s 41ms/step

- binary_accuracy: 0.8150 - loss: 1.3822 - val_binary_accuracy: 0.8148 -
val_loss: 0.4519

Epoch 10/10

7602/7602 299s 39ms/step

- binary_accuracy: 0.8150 - loss: 1.3900 - val_binary_accuracy: 0.8148 -
val_loss: 0.4510


```
[58]: <keras.src.callbacks.history.History at 0x7d7f36bfa350>
```

```
[60]: import matplotlib.pyplot as plt

history = best_model.history

plt.figure(figsize=(10, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['binary_accuracy'])
plt.plot(history.history['val_binary_accuracy'])
plt.title('Model accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')

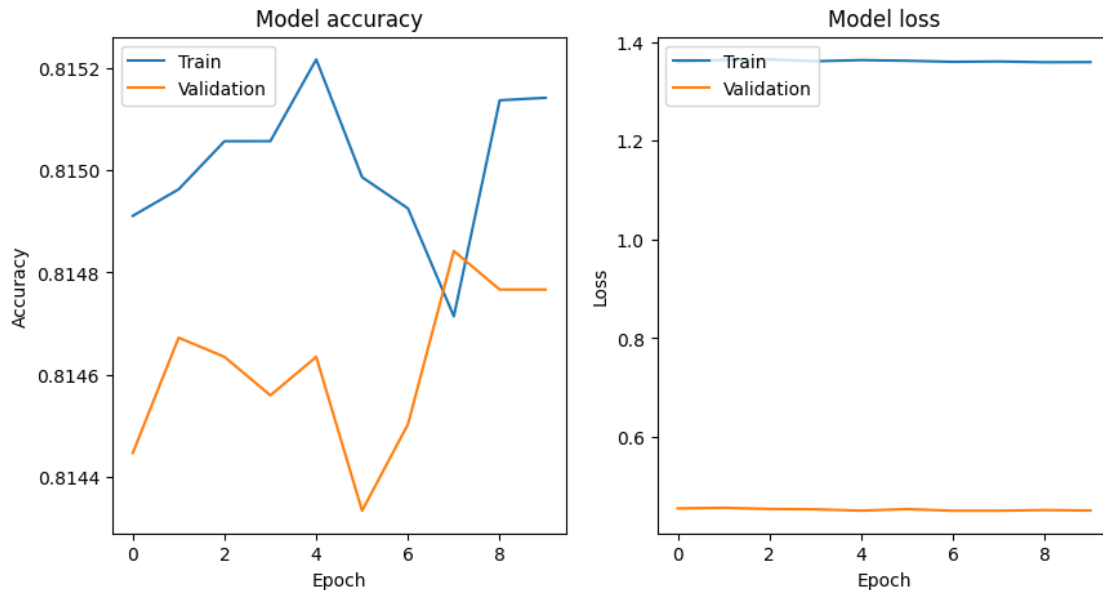
# Plot training & validation loss values
plt.subplot(1, 2, 2)
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('Model loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()

# Example of plotting other metrics (if available in history.history):
if 'precision' in history.history:
    plt.plot(history.history['precision'])
    plt.plot(history.history['val_precision'])
    plt.title('Model Precision')
    plt.ylabel('Precision')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Validation'], loc='upper left')
    plt.show()

if 'recall' in history.history:
    plt.plot(history.history['recall'])
    plt.plot(history.history['val_recall'])
    plt.title('Model Recall')
    plt.ylabel('Recall')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Validation'], loc='upper left')
    plt.show()

if 'f1_score' in history.history:
    plt.plot(history.history['f1_score'])
    plt.plot(history.history['val_f1_score'])
```

```
plt.title('Model F1 Score')
plt.ylabel('F1 Score')
plt.xlabel('Epoch')
plt.legend(['Train', 'Validation'], loc='upper left')
plt.show()
```



```
[63]: # Save the best model to Google Drive
import os

DRIVE_MODEL_PATH = DRIVE_PATH + '/best_model-1/' # Define the path in Google_
↳ Drive, adding .keras extension

# Create the directory if it doesn't exist
!mkdir -p "{DRIVE_MODEL_PATH}"

# Save the model
best_model.save(DRIVE_MODEL_PATH + 'best_model.keras')

print(f"Model saved to: {DRIVE_MODEL_PATH}")

# Verification (optional): Check if the model directory exists and is not empty
if os.path.exists(DRIVE_MODEL_PATH) and os.listdir(DRIVE_MODEL_PATH):
    print("Model directory exists and contains files.")
else:
    print("Error: Model directory does not exist or is empty.")
```

Model saved to: /content/drive/MyDrive/AAI-590_Collabs/best_model-1/

Model directory exists and contains files.